

Gradient boosting with XGBoost

- State of the art data science library for performing classification and regression.
- XGBoost makes use of gradient boosted decision trees, which provide better performance.
- The problem with regular decision is that they are often not complex enough to capture intricacies of many large datasets. We could increase the max depth but it causes overfitting of data.

Instead, we use gradient boosting to combine many decision trees into a single model for classification or regression. Gradient boosting starts with a single decision tree, then iteratively adds more decision trees to overall model.

- The basic structure for XGBoost is DMatrix, which represents a data matrix. DMatrix can be constructed from numpy array.

```
import xgboost as xgb
dmat1 = xgb.DMatrix(data)
dmat2 = xgb.DMatrix(data, label=labels)
```

↑
numpy array.

The DMatrix object can be used to train a Booster object, which represents the gradient boosted decision tree.

```
dtrain = xgb.DMatrix(data, label = labels)
params = {
    'max_depth': 0,
    'objective': 'binary:logistic'
```

```
bst = xgb.train(params, dttrain) # booster
```

```
print(bst.eval(deval)) # evaluation
dpred = xgb.DMatrix(new_data) # new data
predictions = bst.predict(dpred)
print(predictions)
```

Note - predictions from predict function are probabilities and not class labels

Cross-validation

We can tune the parameters using cross validation

```
dttrain = xgb.DMatrix(data, label = labels)
params = {
    'max_depth': 2,
    'lambda': 1.5,
    'objective': 'binary:logistic'
}
cv_result = xgb.cv(params, dttrain)
print(cv_result)
```

The output of cv is a pandas dataframe
It contains the training and testing results
(mean & SD) of a k-fold cross validation.
The value of k is set with the nfold argument.
The keyword num-boost-round specifies the
number of boosting iterations.

→ `cv_results = ngb.cv(params, d_train,
num_boost_round=5)`

Storing Boosters

→ ~~best~~ Storing into a binary file called model.bin
`bst = ngb.train(params, d_train)
bst.save_model('model.bin')`

→ loading a booster
This requires us to create an empty booster
then load the file's data into it.
`new_bst = ngb.Booster()
new_bst = load_model('model.bin')`

XGBoost Classifier

XGBoost classifier takes Numpy arrays as
input arguments

→ `model = ngb.XGBClassifier()
model.fit(data, labels)
predictions = model.predict(new_data)`

Note:- predict function returns classes and not probabilities

- We can set classification type to XGBClassifier Object like in the Original Booster Object.

XGBoost Regressor

XGBoost also provides a scikit-learn style linear regression model.

- `model = xgb.XGBRegressor(max_depth=2)`
`model.fit(data)`

`predictions = model.predict(data_new)`
`print(predictions)`

Feature Importance

- Not every feature in a dataset is used equally for helping a boosted decision tree make a decision.

After training an XGBoost model, we can view the relative (proportional) importance of each dataset feature.

- `print(model.feature_importances_)`

- We can also plot this
`xgb.plot_importance(model)`
`plt.show()` #matplotlib plot

The resulting plot is a bar graph of the F-scores (F₁-scores) for each feature.

We can ~~also~~ pass `importance-type="gain"` to `xgb.plot_importance`. By default it uses `weight`. `gain` → information gain. Information gain is a commonly used metric for determining how good a feature is at differentiating the dataset.

Hyperparameter Tuning

Apply grid search cross validation to XGBoost model

```
→ model = xgb.XGBClassifier()  
params = {'max_depth': range(2, 5)}
```

```
from sklearn.model_selection import GridSearchCV  
cv_model = GridSearchCV(model, params,  
                          cv=4, iid=False)
```

```
cv_model.fit(data, labels)  
print(cv_model.best_params_)
```

Note - A couple of commonly tuned params are `max_depth` and `eta` (learning rate of boosting algorithm)

Model Persistence

Same as `sklearn` using the `joblib` API. Use the `dump` and `load` functions.